Prasanna Kailash Warad
469-766-7241 | waradprasanna@gmail.com | Dallas, TX | http://www.linkedin.com/in/prasannawarad

---

TECHNICAL SKILLS
Data Engineering: ETL, ELT, Airflow, Spark, Hadoop, Data Modeling, Lambda/Kappa Architecture, dbt
Programming: Python (Pandas, NumPy, Matplotlib, Seaborn, PySpark, Scikit-learn, PyTorch), SQL (Advanced)
Databases & Platforms: MySQL, PostgreSQL, SQL Server, Oracle, MongoDB, Snowflake, AWS Redshift, Azure Synapse
Web & Frontend: React, Node.js, REST APIs, CSS
Tools & Visualization: Tableau, Power BI, Figma, Docker, Git, GitHub

EDUCATION
The University of Texas at Dallas, Dallas, TX — Aug 2024 - May 2026
Master of Science, Information Technology & Management, *Dean's Excellence Scholar*
Graduate Certificate in Business Analytics & Data Mining
Savitribai Phule Pune University, Pune, India — Aug 2018 - May 2022
Bachelor of Engineering, Electrical

PROFESSIONAL EXPERIENCE
Cloud BC Labs, Reston, Virginia — Oct 2025 - Present
Software Engineer Intern
- **Developed** and deployed RESTful APIs using **Node.js** and **Python**, streamlining data exchange between front-end and backend services and reducing integration effort by 20%.
- **Restructured** ETL pipelines in **Python** through indexing, query refactoring, and batch processing strategies, improving pipeline throughput by 25% while tuning underlying **SQL** queries for efficiency.
- **Identified** behavioral trends from application usage metrics, cutting recurring bug reports by 15% and accelerating sprint prioritization decisions.

HCLTech Ltd, Noida, India — Feb 2023 - Jul 2024
Data Engineer — Client: Dollar General (Fortune 100 Retailer)
- **Analyzed** 2M+ retail transaction records using **Python** and **SQL** to detect data anomalies and inconsistencies, improving accuracy across operational reporting workflows and strengthening downstream model reliability.
- **Designed** automated **Spark** and **Airflow** data pipelines for real-time operational reporting, reducing query latency by 28% through caching and partition pruning strategies.
- **Built** a predictive model using **scikit-learn** and **PyTorch** on **AWS** system performance logs to detect anomalies and forecast potential downtimes, cutting incident response time by 25%.

PROJECTS
Credit Risk Modeling & Default Prediction Pipeline — Sep 2025 - Dec 2025
- **Constructed** a credit risk pipeline in **Python** and **scikit-learn** on 255K+ lending records, engineering 6 features and applying **SMOTE** to correct a 1:7.6 class imbalance for robust minority-class training.
- **Evaluated** Logistic Regression, Random Forest, and XGBoost via 5-fold CV; XGBoost achieved best performance (88.5% accuracy, ROC-AUC: 0.737) - proposed a hybrid deployment strategy balancing precision and recall.

Scalable Fleet Risk Analytics Pipeline — Feb 2025 - Apr 2025
- **Architected** a scalable data pipeline using **Hadoop (HDFS, MapReduce)** to ingest and transform operational datasets across 300+ assets in 14 states, enabling compliance monitoring and performance tracking at scale.
- **Implemented** data ingestion and storage strategies to support efficient querying and integration with **Power BI** dashboards for asset risk monitoring and predictive maintenance reporting.

Instacart Customer Behavior Dashboard — Jan 2025 - Feb 2025
- **Processed** 3M+ grocery transaction records through optimized **SQL** pipelines to support operational reporting for customer segmentation, cohort analysis, and demand forecasting.
- **Integrated** cleaned datasets into **Tableau** dashboards tracking operational KPIs — product demand, fulfillment trends, and campaign ROI — enabling data-driven decisions for cross-functional stakeholders.